

# Algorithms overview

## Pre-strike

- ▶ Prüfer code
- ▶ Cheapest Spanning Trees: Kruskal + Prim

## What's left:

- ▶ Dijkstra's Algorithm for Shortest Path
- ▶ Traveling Salesperson Problem


## What we're going to skip:

Longest path + applications to scheduling (3.5 in notes)

# Directed graphs

## Definition

A *directed graph* is one where each edge has a chosen starting and ending point, usually indicated with arrows.



## Walks in directed graphs:

You can only travel the edge in the direction the arrow shows.

# Dijkstra's algorithm for shortest path

## Input:

A weighted (possibly directed) graph  $G$  and starting vertex  $v \in G$

## Output:

For every vertex  $w \in G$  a list of all shortest paths from  $v$  to  $w$

## Initialize:


From starting vertex  $v$  list every edge out of  $v$  as a potential shortest path to corresponding vertex  $w$

## Iterate:

- ▶ Choose  $w$  with cheapest potential shortest path and make these paths permanent
- ▶ Update list of potential paths by adding edges out of  $w$  to the shortest paths to  $w$  and checking if they're cheaper than known paths

## Example graph from the 2008 Exam

Find all shortest paths from  $S$  to  $T$ .



Add on:

- ▶ Which edges if made a little longer would make the distance from  $S$  to  $T$  longer?
- ▶ Which edges if made a little shorter would make the distance from  $S$  to  $T$  shorter?

# The main idea of why Dijkstra's works:

Suppose we're at the step where Dijkstra's decides the cheapest path to  $w$ .

Then:

A cheaper path to  $w$  would have to go through a vertex  $u$  we haven't found the cheapest path to yet.

But!

Even *getting* to  $u$  costs more than our cheapest path to  $w$ .

An observation:

Dijkstra's algorithm depends on the edge weights being non-negative!

# Culture: performance of Dijkstra's algorithm

In a limited sense, Dijkstra's algorithm is optimal:

If *all we know* is that we have a weighted graph, then you can't do better than Dijkstra's algorithm.

In practise, often not very good:

When finding path from Sheffield to Edinburgh, Dijkstra's algorithm explores every street in London.

Real world maps have extra information:

It's easy to calculate the distance between two points as the crow flies, and we know the driving distance has to be at least that large.

The A\* algorithm avoids searching London:

Supposes we have an easy to calculate "heuristic distance"  $h(v, w)$ , that is a lower bound for the actual distance  $d(v, w)$ .